



## Operating Systems

# “Threads, SMP, and Microkernels”



## Roadmap



- ➔ • Threads: Resource ownership and execution
  - Symmetric multiprocessing (SMP).
  - Microkernel
  - Case Studies of threads and SMP:
    - Windows
    - Solaris
    - Linux



## Processes and Threads

- Processes have two characteristics:
  - **Resource ownership** - process includes a virtual address space to hold the process image
  - **Scheduling/execution** - follows an execution path that may be interleaved with other processes
- These two characteristics are treated independently by the operating system



## Processes and Threads

- The dispatching unit is referred to as a **thread** or lightweight process
- The entity that owns a resource is referred to as a process or **task**
- One process / task can have one or more threads

# Multithreading

- The ability of an OS to support multiple, concurrent paths of execution within a single process.




Figure 4.1 Threads and Processes [ANDE97]

## Single Thread Approaches

- MS-DOS supports a single user process and a single thread.
- Some UNIX, support multiple user processes but only support one thread per process




Figure 4.1 Threads and Processes [ANDE97]

# Multithreading

- Java run-time environment is a single process with multiple threads
  - Multiple processes **and** threads are found in Windows, Solaris, and many modern versions of UNIX
- the main topic




Figure 4.1 Threads and Processes [ANDE97]

# Processes

- A virtual address space which holds the process image
- Protected access to
  - Processors,
  - Other processes,
  - Files,
  - I/O resources



## One or More Threads in Process

- Each thread has
  - An execution state (running, ready, etc.)
  - Saved thread context when not running
  - An execution stack
  - Some per-thread static storage for local variables
  - Access to the memory and resources of its process (all threads of a process share this)



## One view...

- *One way to view a thread is as an independent program counter operating within a process.*

## Threads vs. processes

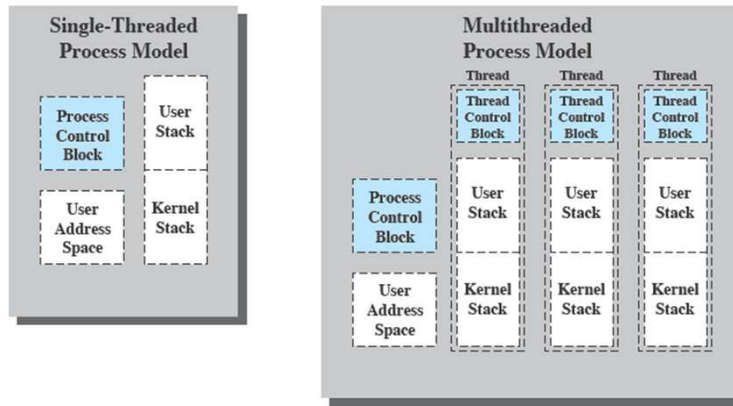


Figure 4.2 Single Threaded and Multithreaded Process Models

## Benefits of Threads

- Takes less time to create a new thread than a process
- Less time to terminate a thread than a process
- Switching between two threads takes less time than switching processes
- Threads can communicate with each other
  - without invoking the kernel



## Thread use in a Single-User System

- Foreground and background work
- Asynchronous processing
- Speed of execution
- Modular program structure



## Threads

- Several actions that affect all of the threads in a process
  - The OS must manage these at the process level.
- Examples:
  - Suspending a process involves suspending all threads of the process
  - Termination of a process, terminates all threads within the process



## Activities similar to Processes

- Threads have execution states and may synchronize with one another.
  - Similar to processes
- We look at these two aspects of thread functionality in turn.
  - States
  - Synchronisation



## Thread Execution States


- States associated with a change in thread state
  - Spawn (another thread)
  - Block
    - Issue: will blocking a thread block other, or all threads
  - Unblock
  - Finish (thread)
    - Deallocate register context and stacks



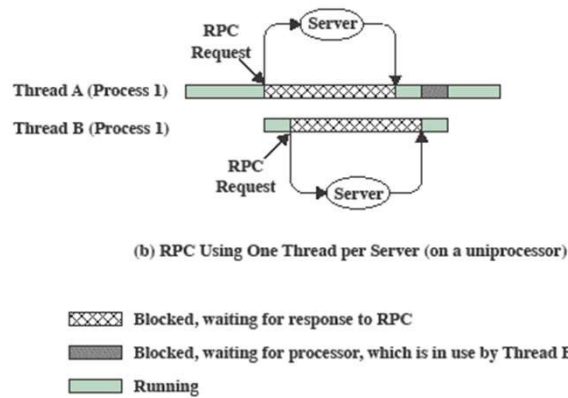
## Example: Remote Procedure Call

- Consider:
  - A program that performs two remote procedure calls (RPCs)
  - to two different hosts
  - to obtain a combined result.

## RPC Using Single Thread



## RPC Using One Thread per Server



## Multithreading on a Uniprocessor

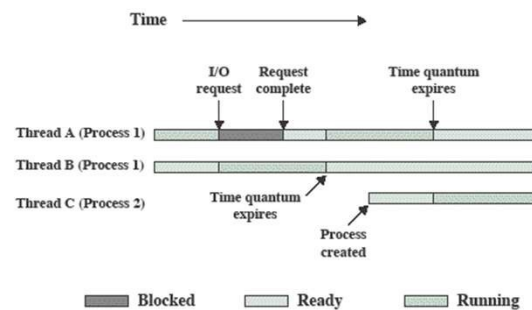



Figure 4.4 Multithreading Example on a Uniprocessor

## Categories of Thread Implementation


- User Level Thread (ULT)
  - thread that are managed entirely by user-level code
  - Not requiring any support from the operating system kernel.
- Kernel level Thread (KLT) also called:
  - kernel-supported threads
  - lightweight processes.

## User-Level Threads

- All thread management is done by the application
- The kernel is not aware of the existence of threads



## Kernel-Level Threads



(b) Pure kernel-level

- Kernel maintains context information for the process and the threads
  - No thread management done by application
- Scheduling is done on a thread basis
- Windows is an example of this approach

## Advantages of KLT

- The kernel can simultaneously schedule multiple threads from the same process on multiple processors.
- If one thread in a process is blocked, the kernel can schedule another thread of the same process.
- Kernel routines themselves can be multithreaded.

## Disadvantage of KLT


- The transfer of control from one thread to another within the same process requires a mode switch to the kernel

## Comparison ULT and KLT

Criteria	User-Level Threads (ULTs)	Kernel-Level Threads (KLTs)
Management	Managed entirely by user-level code, without kernel support	Managed by the operating system kernel
Scheduling	Scheduled by a user-level thread library or application, without kernel intervention	Scheduled by the operating system kernel, which may provide advanced scheduling algorithms and policies
Context Switching	Context switching occurs entirely in user space, without requiring a system call or trap into kernel mode	Context switching requires a system call or trap into kernel mode, which can incur overhead
Resources	ULTs share the same process address space and resources as the parent process	KLTs have their own kernel-level data structures, which can result in higher overhead and memory usage
Scalability	ULTs are limited to a single processor core and cannot take full advantage of multicore systems	KLTs can be assigned to different processor cores and can take full advantage of multicore systems
Synchronization	ULTs must use user-level synchronization mechanisms, which can be more efficient but may be subject to priority inversion and other issues	KLTs can use both user-level and kernel-level synchronization mechanisms, providing more flexibility and better enforcement of policies
Portability	ULTs are highly portable across different operating systems, as long as a compatible user-level thread library is available	KLTs may be less portable, as different operating systems may have different thread APIs and scheduling mechanisms

## Combined Approaches

- Thread creation done in the user space
- Bulk of scheduling and synchronization of threads by the application
- Example is Solaris



(c) Combined

## Relationship Between Thread and Processes

Table 4.2 Relationship Between Threads and Processes

Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	Combines attributes of M:1 and 1:M cases.	TRIX



## Roadmap

- Threads: Resource ownership and execution
- • Symmetric multiprocessing (SMP).
- Microkernel
- Case Studies of threads and SMP:
  - Windows
  - Solaris
  - Linux



## Traditional View

- Traditionally, the computer has been viewed as a sequential machine.
  - A processor executes instructions one at a time in sequence
  - Each instruction is a sequence of operations
- Two popular approaches to providing parallelism
  - Symmetric Multi Processors (SMPs)
  - Clusters (ch 16)



## Categories of Computer Systems

- Single Instruction Single Data (SISD) stream
  - Single processor executes a single instruction stream to operate on data stored in a single memory
- Single Instruction Multiple Data (SIMD) stream
  - Each instruction is executed on a different set of data by the different processors



## Categories of Computer Systems

- Multiple Instruction Single Data (MISD) stream
  - A sequence of data is transmitted to a set of processors, each of execute a different instruction sequence
- Multiple Instruction Multiple Data (MIMD)
  - A set of processors simultaneously execute different instruction sequences on different data sets



## Parallel Processor Architectures

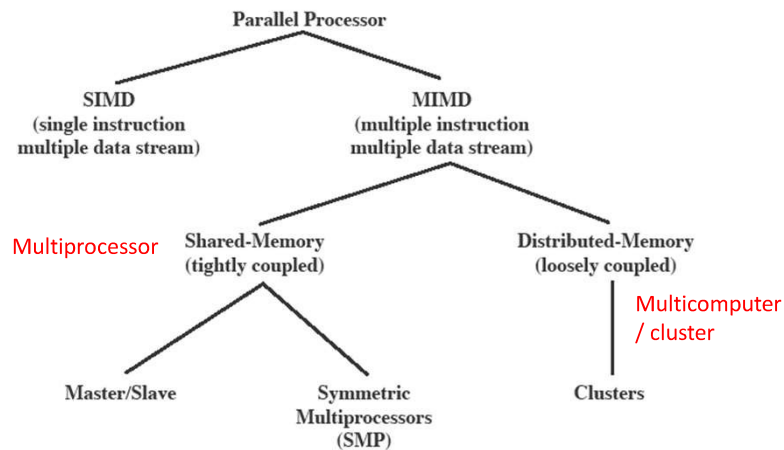


Figure 4.8 Parallel Processor Architectures

## Symmetric Multiprocessing

- Kernel can execute on any processor
  - Allowing portions of the kernel to execute in parallel
- Typically, each processor does self-scheduling from the pool of available process or threads

## Typical SMP Organization

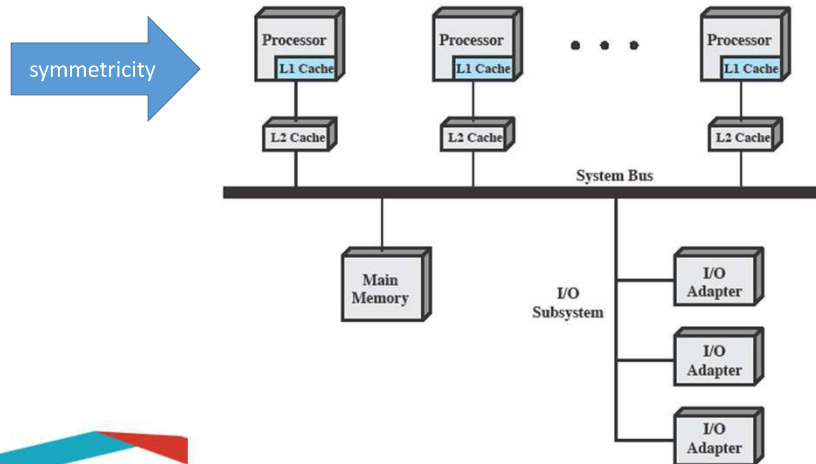


Figure 4.9 Symmetric Multiprocessor Organization

## Multiprocessor OS Design Considerations



- The key design issues include
  - Simultaneous concurrent processes or threads
  - Scheduling
  - Synchronization
  - Memory Management
  - Reliability and Fault Tolerance



## Roadmap

- Threads: Resource ownership and execution
- Symmetric multiprocessing (SMP).
- • **Microkernel**
- Case Studies of threads and SMP:
  - Windows
  - Solaris
  - Linux



## Microkernel

- A microkernel is a small OS core that provides the foundation for modular extensions.
- Big question is how small must a kernel be to qualify as a microkernel
  - *Must* drivers be in user space?
- In theory, this approach provides a high degree of flexibility and modularity.