Operating Systems:
Internals and Design Principles, 9/E
William Stallings

# Chapter 7
# Memory Management

Patricia Roy
Manatee Community College, Venice, FL


# Why?

*Memory needs to be allocated to ensure a reasonable supply of ready processes to consume available processor time*

# Roadmap

- Basic requirements of Memory Management
- Memory Partitioning
- Basic blocks of memory management
  – Paging
  – Segmentation

# The need for memory management

- Memory is cheap today, and getting cheaper
  – But applications are demanding more and more memory, there is never enough!
- Memory Management, involves swapping blocks of data from secondary storage.
- Memory I/O is slow compared to a CPU
  – The OS must cleverly time the swapping to maximise the CPU's efficiency

# Memory Management scope

- Relocation
- Protection
- Sharing
- Logical organisation
- Physical organisation

# Relocation

- The programmer does not know where the program will be placed in memory when it is executed,
  - it may be swapped to disk and return to main memory at a different location (relocated)
- Memory references must be translated to the actual physical memory address

# Memory Management Terms

**Table 7.1 Memory Management Terms**

| Term | Description |
|---|---|
| Frame | ***Fixed***-length block of main memory. |
| Page | ***Fixed***-length block of data in secondary memory (e.g. on disk). |
| Segment | ***Variable-length*** block of data that resides in secondary memory. |

# Addressing

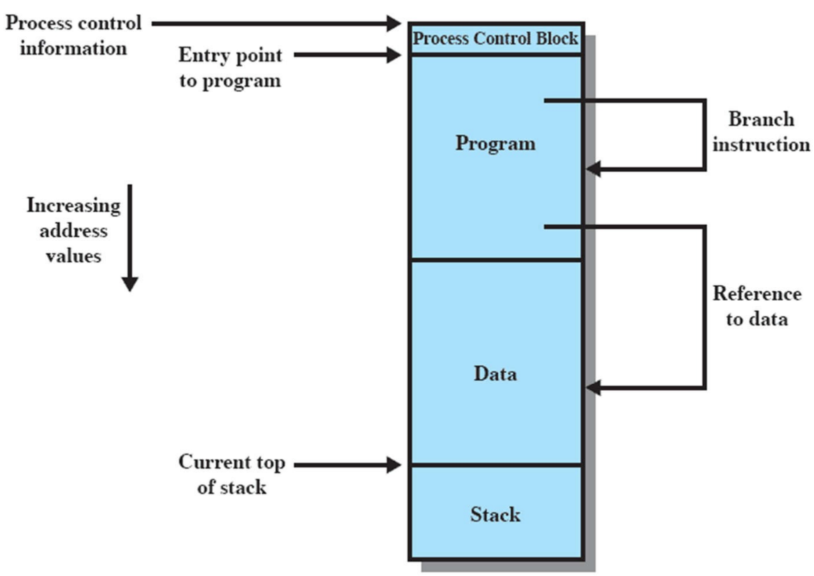


**Figure 7.1 Addressing Requirements for a Process**

# Protection

- Processes should not be able to reference memory locations in another process without permission
- Impossible to check absolute addresses at compile time
- Must be checked at run time

# Sharing

- Allow several processes to access the same portion of memory
- Better to allow each process access to the same copy of the program rather than have their own separate copy

# Logical Organization

- Memory is organized linearly (usually)
- Programs are written in modules
  - Modules can be written and compiled independently
- Different degrees of protection given to modules (read-only, execute-only)
- Share modules among processes

# Physical Organization

- Cannot leave the programmer with the responsibility to manage memory
- Memory available for a program plus its data may be insufficient
  - Overlaying allows various modules to be assigned the same region of memory but is time consuming to program
- Programmer does not know how much space will be available

# Partitioning

- An early method of managing memory
- divides primary **memory** into multiple **memory partitions**, usually contiguous areas of **memory**
- Each **partition** might contain all the information for a specific job or task.
- **Memory management** consists of allocating a **partition** to a job when it starts and unallocating it when the job ends.

# Types of Partitioning

- Fixed Partitioning
- Dynamic Partitioning
- Simple Paging
- Simple Segmentation
- Virtual Memory Paging
- Virtual Memory Segmentation

# Fixed Partitioning

- Equal-size partitions (see fig 7.3a)
  - Any process whose size is less than or equal to the partition size can be loaded into an available partition
- The operating system can swap a process out of a partition
  - If none are in a ready or running state

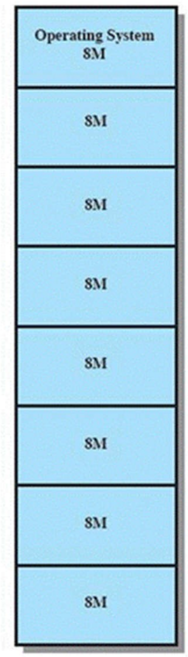


(a) Equal-size partitions

# Fixed Partitioning Problems

- A program may not fit in a partition.
  - The programmer must design the program with overlays
- Main memory use is inefficient.
  - Any program, no matter how small, occupies an entire partition.
  - This is results in ***internal fragmentation.***

# Solution – Unequal Size Partitions

- Lessens both problems
  – but doesn't solve completely
- In Fig 7.3b,
  – Programs up to 16M can be accommodated without overlay
  – Smaller programs can be placed in smaller partitions, reducing internal fragmentation
- Still a possible internal fragmentation problem

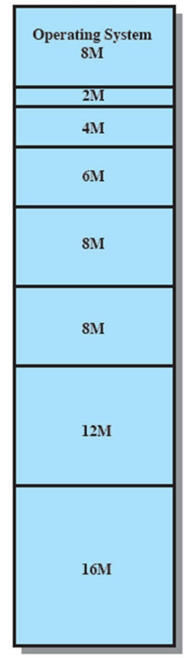


(b) Unequal-size partitions

# Placement Algorithm

- Equal-size
  – Placement is trivial (no options)
- Unequal-size
  – Can assign each process to the smallest partition within which it will fit
  – Queue for each partition
  – Processes are assigned in such a way as to minimize wasted memory within a partition

# Fixed Partitioning

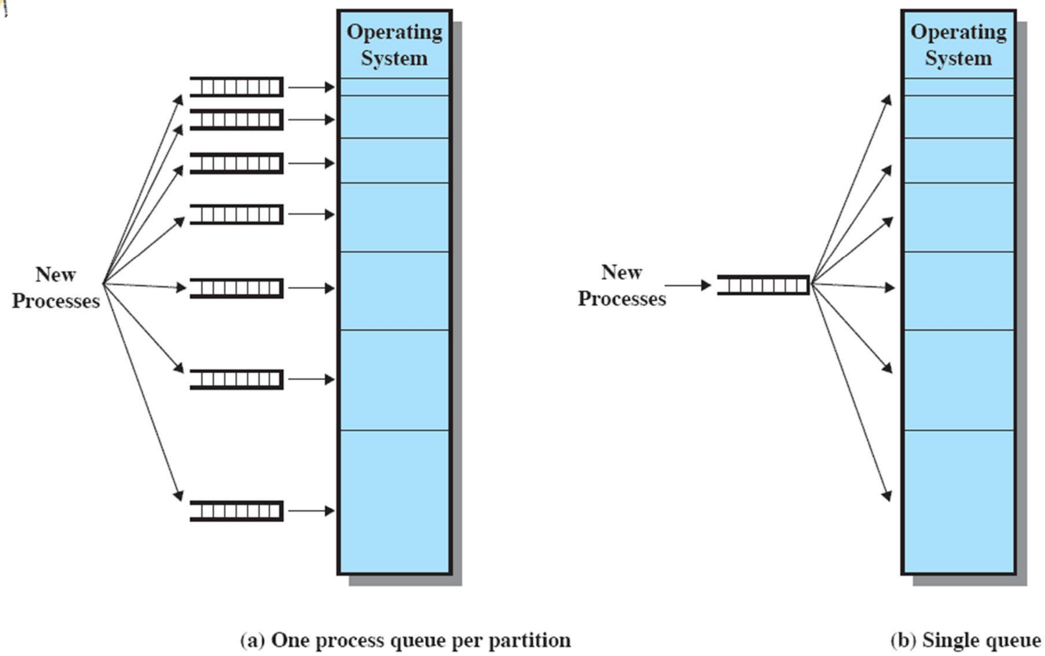


Figure 7.3 Memory Assignment for Fixed Partitioning

# Remaining Problems with Fixed Partitions

- The number of active processes is limited by the system
  - i.e. limited by the pre-determined number of partitions
- A large number of very small process will not use the space efficiently
  - In either fixed or variable length partition methods

# Dynamic Partitioning

- Partitions are of variable length and number
- Process is allocated exactly as much memory as required

# Dynamic Partitioning Example


Refer to Figure 7.4

- ***External Fragmentation***
- Memory external to all processes is fragmented
- Can resolve using ***compaction***
  - OS moves processes so that they are contiguous
  - Time consuming and wastes CPU time

# Dynamic Partitioning

- Operating system must decide which free block to allocate to a process
- Best-fit algorithm
  – Chooses the block that is closest in size to the request
  – Worst performer overall
  – Since smallest block is found for process, the smallest amount of fragmentation is left
  – Memory compaction must be done more often

# Dynamic Partitioning

- First-fit algorithm
  – Scans memory form the beginning and chooses the first available block that is large enough
  – Fastest
  – May have many process loaded in the front end of memory that must be searched over when trying to find a free block

# Dynamic Partitioning

- Next-fit
  - Scans memory from the location of the last placement
  - More often allocate a block of memory at the end of memory where the largest block is found
  - The largest block of memory is broken up into smaller blocks
  - Compaction is required to obtain a large block at the end of memory

# Allocation

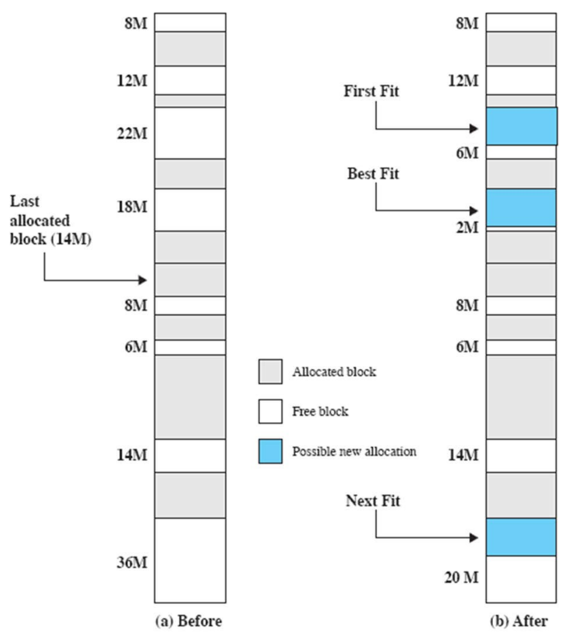


**Figure 7.5** Example Memory Configuration before and after Allocation of 16-Mbyte Block

# Buddy System

- Entire space available is treated as a single block of $2^U$
- If a request of size $s$ where $2^{U-1} < s <= 2^U$
  - entire block is allocated
- Otherwise, block is split into two equal buddies
  - Process continues until smallest block greater than or equal to $s$ is generated

# Example of Buddy System

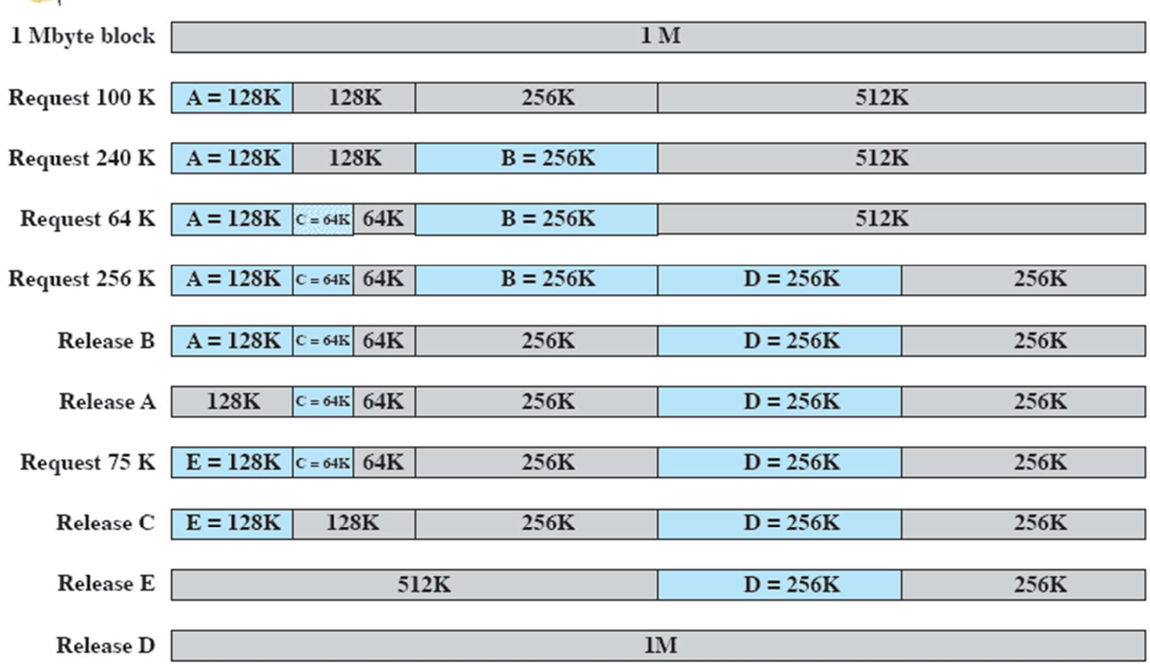


**Figure 7.6 Example of Buddy System**

# Tree Representation of Buddy System

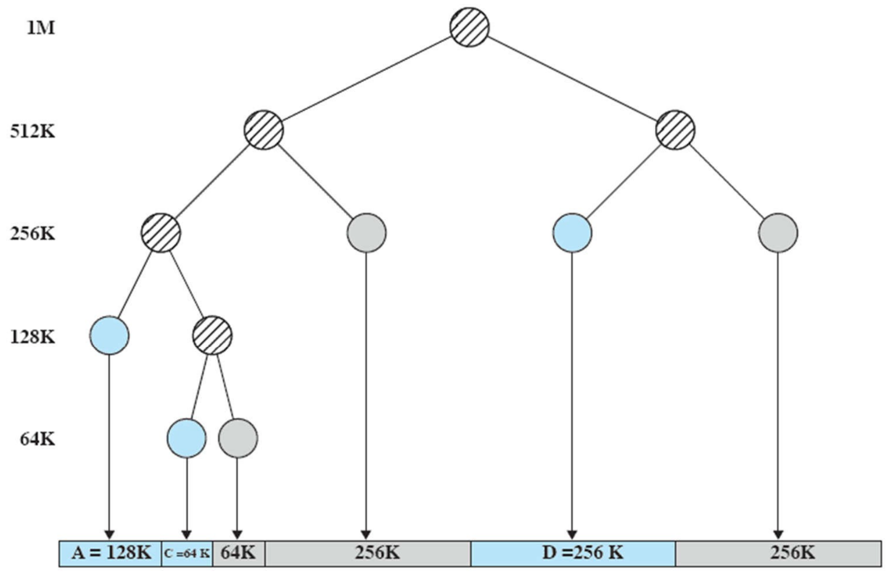


Figure 7.7 Tree Representation of Buddy System

# Relocation

- When program loaded into memory the actual (absolute) memory locations are determined
- A process may occupy different partitions which means different absolute memory locations during execution, due to:
  - Swapping
  - Compaction

# Addresses

- Logical
  - Reference to a memory location independent of the current assignment of data to memory.
- Relative
  - Address expressed as a location relative to some known point.
- Physical or Absolute
  - The absolute address or actual location in main memory.

# Relocation

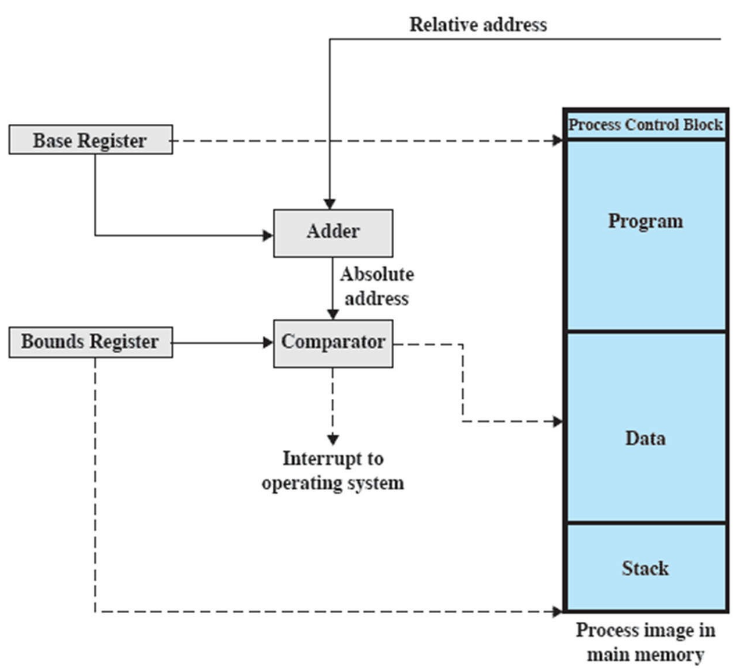


**Figure 7.8 Hardware Support for Relocation**

# Registers Used during Execution

- Base register
  - Starting address for the process
- Bounds register
  - Ending location of the process
- These values are set when the process is loaded or when the process is swapped in

# Registers Used during Execution

- The value of the base register is added to a relative address to produce an absolute address (physical address)
- The resulting address is compared with the value in the bounds (the max address in a memory) register
- If the address is not within bounds, an interrupt is generated to the operating system

# Paging

- Partition memory into small equal fixed-size chunks and divide each process into the same size chunks
- The chunks of a process are called ***pages***
- The chunks of memory are called ***frames***

# Paging

- Operating system maintains a page table for each process
  – Contains the frame location for each page in the process
  – Memory address consist of a page number and offset within the page

# Processes and Frames

| Frame number | Main memory |
|---|---|
| 0 | A.0 |
| 1 | A.1 |
| 2 | A.2 |
| 3 | A.3 |
| 4 | D.0 |
| 5 | D.1 |
| 6 | D.2 |
| 7 | C.0 |
| 8 | C.1 |
| 9 | C.2 |
| 10 | C.3 |
| 11 | D.3 |
| 12 | D.4 |
| 13 | |
| 14 | |

# Page Table

Process A page table

| | |
|---|---|
| 0 | 0 |
| 1 | 1 |
| 2 | 2 |
| 3 | 3 |

Process B page table

| | |
|---|---|
| 0 | — |
| 1 | — |
| 2 | — |

Process C page table

| | |
|---|---|
| 0 | 7 |
| 1 | 8 |
| 2 | 9 |
| 3 | 10 |

Process D page table

| | |
|---|---|
| 0 | 4 |
| 1 | 5 |
| 2 | 6 |
| 3 | 11 |
| 4 | 12 |

Free frame list

| |
|---|
| 13 |
| 14 |

Figure 7.10 Data Structures for the Example of Figure 7.9 at Time Epoch (f)

# Segmentation

- A program can be subdivided into segments
  - Segments may vary in length
  - There is a maximum segment length
- Addressing consist of two parts
  - a segment number and
  - an offset
- Segmentation is similar to dynamic partitioning

# Logical Addresses

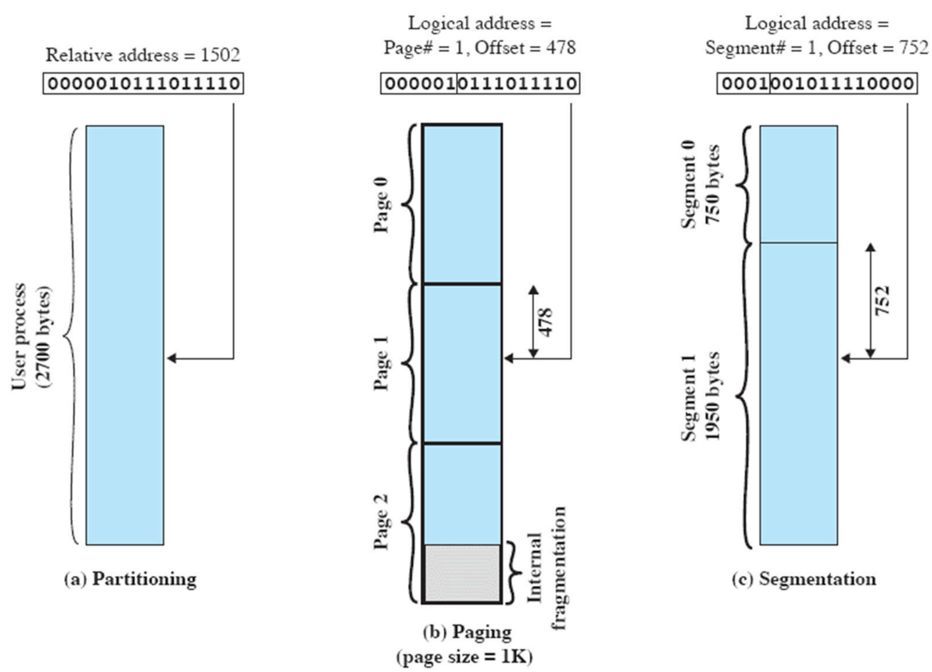


Figure 7.11 Logical Addresses

# Paging

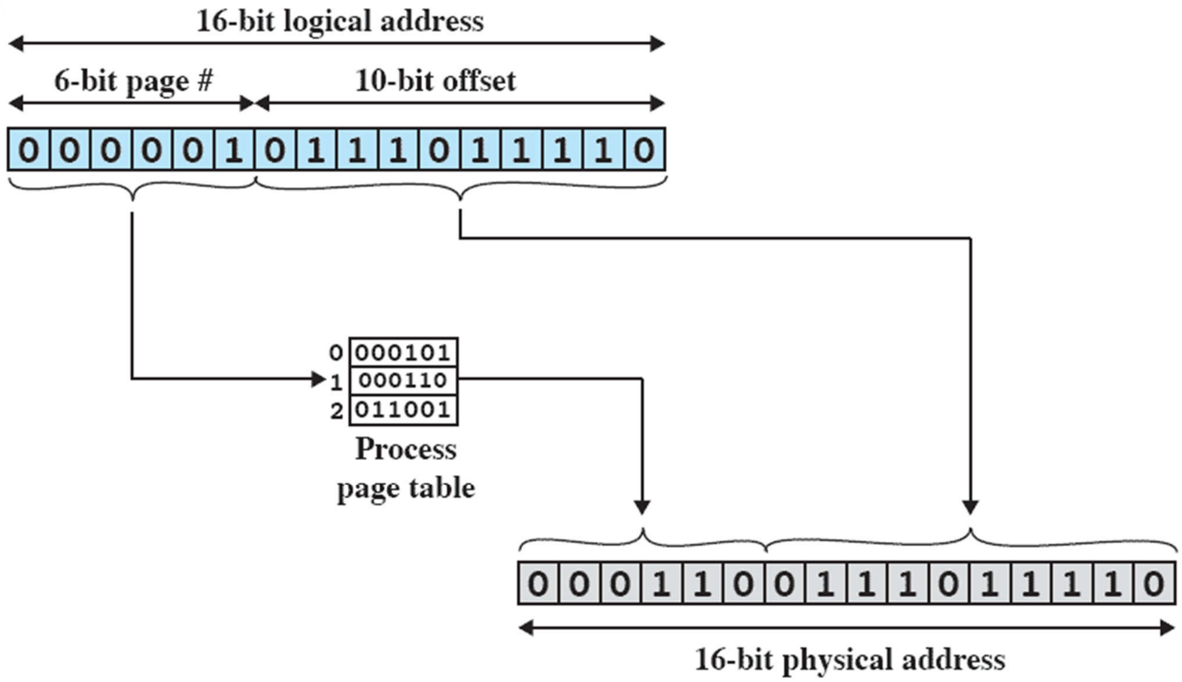

(a) Paging

# Segmentation

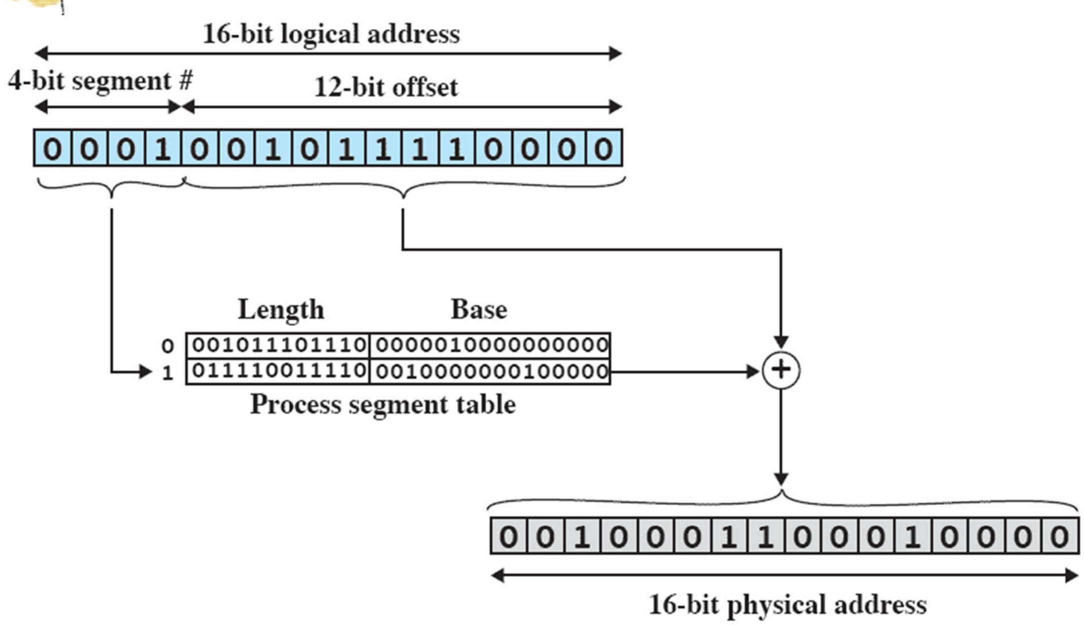

(b) Segmentation

**Figure 7.12 Examples of Logical-to-Physical Address Translation**